


Department of Computer Science and Engineering



SKIN CANCER DETECTION SYSTEM DEEP LEARNING



BACHELOR OF TECHNOLOGY COMPUTER SCIENCE AND ENGINEERING

SUBMITTED TO:

MR SAURAV GUPTA

TRAINER

CSED

SUBMITTED BY:

TEAM NO: 6

COURSE: B.Tech. (CSE), SEC: A

SEMESTER: 5

TEAM LEADER: Sushant Gupta

TOTAL MEMBERS: 10

Invertis University, Bareilly [UP]

PROJECT REPORT (2025-26)

CERTIFICATE

This is to certify that the project report titled “**SKIN CANCER DETECTION SYSTEM**” is the original work of our team, completed in partial fulfillment of the requirements for the degree of **Bachelor of Computer Science and Engineering** at **Invertis University**.

This project has been carried out under the guidance and supervision of **Mr. Saurav Gupta**. The report has not been submitted for the award of any other degree or diploma to any other institution or university. The research presented in this project is a valuable contribution to understanding escalating rates of Cancer and more efficient diagnostic tools. We I hereby certify that the work embodied in this project report is genuine and has been carried out by our team with integrity and dedication.

Mentor: MR. SAURAV GUPTA

Invertis University

ACKNOWLEDGMENT

We would like to express the deepest gratitude to Invertis University, Bareilly for giving us the opportunity to undertake this project titled “**Skin Cancer detection System**” using **Deep Learning**. This project provided us with great learning experience and enhanced our knowledge about integrating Internet of things with artificial intelligence for real-world healthcare applications.

We sincerely thank our department for providing guidance, facilitates and constant encouragement. We are also thankful to our friends for their valuable suggestions and feedback during the development of this project. I am deeply thankful to my *peers and fellow students* for their valuable inputs, collaboration, and shared learning experiences, which contributed significantly to the completion of this project. Their discussions and suggestions provided new perspectives that enriched my understanding of data analytics and consumer research.

Lastly, I extend my sincere gratitude to everyone who directly or indirectly supported me during this project. This experience has been immensely insightful and has strengthened our interest in Deep Learning, its applications in cancer detection and future uses.

We sincerely hope that this project will serve as a useful contribution to the **field of cancer detection** and help in preventing or curing it by detection at early stages.

ABSTRACT

Project Title: Skin Cancer Detection System

Domain: Deep Learning, Cancer Detection

Technology Stack: TensorFlow, CNN, Data Visualization

The **Skin Cancer Detection using Deep Learning** project aims to develop an intelligent deep learning-based system capable of accurately identifying various types of skin cancer from dermoscopic images. Leveraging the power of the **artificial intelligence**, this project integrates real-time data processing, cloud-based computation, and deep **convolutional neural networks (CNNs)** to assist in early diagnosis and classification of skin cancer types.

The system utilizes the **ISIC Skin Cancer Dataset**, which contains thousands of labeled images across multiple classes such as melanoma, basal cell carcinoma, and benign keratosis. The dataset is imported from **Kaggle** using the kagglehub library and processed using **TensorFlow**, **NumPy**, **Pandas**, and **Matplotlib** for visualization and analysis.

A deep learning model based on **TensorFlow/Keras** is trained to classify skin lesions into distinct categories. The dataset is divided into training, validation, and testing sets, ensuring balanced representation of all classes. Data augmentation techniques are applied to enhance model generalization and handle class imbalance. The training pipeline includes convolutional, pooling, dropout, and fully connected layers to extract relevant image features and make predictions.

Once trained, the model's performance is evaluated using accuracy metrics, confusion matrices, and visualization plots of training curves. The results demonstrate the capability of CNN-based architectures to achieve high precision in medical image classification tasks.

In an IIoT context, this model can be deployed on interconnected devices—such as smart diagnostic systems or hospital servers—enabling automated, remote, and real-time skin cancer screening. This integration reduces manual diagnostic workload, improves detection speed, and enhances healthcare accessibility in remote areas.

Keywords: Skin Cancer Detection using Deep Learning, artificial intelligence, convolutional neural networks, ISIC Skin cancer dataset, kaggle, tensorflow, numpy, pandas, matplotlib, tensorflow/keras

TEAM MEMBERS

S. no	Student Name	Student Id	Contribution
1	Shruti Mishra	BCS2023270	100%
2	Anshika Srivastava	BCS2023272	100%
3	Nitin Singh	BCS2023278	100%
4	Preet Gupta	BCS2023282	100%
5	Atharva Srivastava	BCS2023285	100%
6	Sushant Gupta	BCS2023287	100%
7	Divyanka Babbar	BCS2023294	100%
8	Rishika Maheshwari	BCS2023354	100%
9	Ishita Singh	BCS2023362	100%
10	Khushnam Khatoon	BCS2023375	100%

TABLE OF CONTENTS

CHAPTER 1: INTRODUCTION

1.1 Background	
1.2 Motivation	
1.3 Problem Statement	
1.4 Objectives	
1.5 Scope of the Project	
1.6 Project Overview	

CHAPTER 2: LITERATURE REVIEW

2.1 System Analysis	
2.2 Research Gap	

CHAPTER 3: METHODOLOGY

3.1 Data Collection	
3.2 Data Cleaning & Preprocessing	
3.3 Data Modeling.....	

CHAPTER 4: SYSTEM DESIGN

4.1 Tools & Technologies Used	
-------------------------------------	--

CHAPTER 5: RESULTS & ANALYSIS

5.1 Insights / Observations	
-----------------------------------	--

CHAPTER 6: CONCLUSION & FUTURE SCOPE

6.1 Conclusion	
6.2 Future Scope	

CHAPTER 1: INTRODUCTION

1.1 Background

Skin cancer is one of the most common and rapidly increasing forms of cancer worldwide. According to the World Health Organization (WHO), millions of new skin cancer cases are diagnosed each year, with a significant portion resulting in fatalities due to late detection. Traditionally, diagnosis relies on dermatologist expertise and biopsy-based confirmation, which are both time-consuming and resource-intensive. In many parts of the world, especially rural and remote areas, access to specialized dermatological care is limited, leading to delays in diagnosis and treatment. Recent advancements in Artificial Intelligence (AI) and Deep Learning (DL) have demonstrated remarkable success in medical image analysis. Convolutional Neural Networks (CNNs) are particularly effective in identifying patterns and visual features within images—making them highly suitable for diagnosing diseases like skin cancer from dermoscopic images.

1.2 Motivation

The primary motivation behind this project is the need for early and accurate skin cancer detection using advanced technologies that can be deployed in both clinical and remote environments. Early detection significantly increases survival rates and reduces the cost and complexity of treatment. However, manual diagnosis depends heavily on expert interpretation, which may vary between doctors and regions. This subjectivity, combined with limited availability of dermatologists, creates a bottleneck in efficient diagnosis. This project is driven by the vision of developing a cost-effective, scalable, and intelligent skin cancer detection system that leverages AI and IIoT to enhance healthcare delivery and accessibility worldwide.

1.3 Problem Statement

In many developing and rural regions, access to dermatologists and advanced diagnostic tools is limited. Manual examination of skin lesions is prone to human error and interpretation bias, leading to inconsistent results. With the growing volume of medical imaging data, traditional diagnostic methods cannot scale efficiently. Late-stage detection drastically reduces survival chances, highlighting the need for faster, automated systems. To address these challenges, this project proposes the development of an AI-powered skin cancer detection system integrated within an IIoT framework, capable of classifying skin cancer images into multiple categories using deep learning models (CNNs). The system aims to provide accurate, automated, and accessible diagnostic support for healthcare professionals and remote patients alike.

1.4 Objectives

The core objective is to create an automated and highly accurate diagnostic aid for skin cancer. This is achieved by developing a machine learning model capable of classifying dermatoscopic images of skin lesions. The goal is to provide a non-invasive, rapid, and objective preliminary analysis, thereby improving the efficiency and consistency of early skin cancer detection, which is crucial for patient outcomes.

1.5 Scope of the Project

The project's scope is centered on a detailed, multi-class image classification task using a large, publicly available dataset. The system uses the HAM10000 dataset for training, validation, and testing. The primary methodology involves leveraging a convolutional neural network (CNN) architecture, specifically employing transfer learning with the MobileNetV2 model. This project is framed within an IIoT context, implying that the model is designed not just for high accuracy but also for efficient computation suitable for deployment in a connected, potentially resource-constrained healthcare environment. This suggests the eventual application involves the model operating on an edge device (like a smart dermatoscope or a local gateway) for real-time analysis, integrated into a broader network of medical devices.

1.6 Project Overview

The project successfully implements a deep learning solution to a critical medical diagnostic problem. The system's methodology is centered on the application of a Convolutional Neural Network (CNN). Specifically, it employs the robust technique of transfer learning by fine-tuning the MobileNetV2 architecture. MobileNetV2 is selected for its efficiency, which is crucial for its intended application within an Industrial Internet of Things (IIoT) framework, suggesting the final model is optimized for deployment on low-power, connected edge devices for real-time analysis at the point of care. The implementation involves standard deep learning steps—data preprocessing, augmentation, and training the fine-tuned model using TensorFlow/Keras—to ensure high diagnostic accuracy and reliability across all nine lesion categories.

CHAPTER 2: LITERATURE REVIEW

2.1 System Analysis

The system analysis for the Skin Cancer Detection project, which employs an Industrial Internet of Things (IIoT) based Convolutional Neural Network (CNN), details the system's architecture, data handling, core processes, and identified constraints.

1. System Architecture and Operational Context

The project is structured as a dedicated diagnostic engine within a larger IIoT healthcare framework.

- **Input Acquisition Layer:** The system begins with an implied sensor or imaging device, such as a connected dermatoscope, responsible for capturing the dermoscopic images of skin lesions.
- **Processing Layer (CNN Model):** The central computational element is the deep learning classifier. It uses the MobileNetV2 architecture, implemented using the TensorFlow/Keras framework. This model choice is strategic, prioritizing computational efficiency suitable for deployment on edge computing devices, which is characteristic of an IIoT system. The training environment leverages a high-performance resource (a T4 GPU).
- **Output and Integration Layer:** The final system output is the predicted classification of the skin lesion. In a production IIoT scenario, this result would be immediately relayed to a clinician, potentially displayed on a local interface, and transmitted to a secure central medical record system.

2. Core Process Methodology (Transfer Learning)

The engine relies on a specific machine learning strategy to achieve high diagnostic accuracy:

- **Model Selection:** The MobileNetV2 base model is chosen for its streamlined CNN structure, which is vital for efficient operation on resource-constrained IIoT devices.
- **Transfer Learning Implementation:** The system utilizes transfer learning, where the foundational layers of MobileNetV2 (already trained on a vast general image dataset like ImageNet) are initially frozen. Only the later layers and a new classification head are trained, or "fine-tuned," using the specialized dermoscopic image data. This rapidly adapts the model to the complex visual patterns of skin lesions.
- **Classification Scope:** The model is tasked with accurately distinguishing between nine categories of lesions: actinic keratosis, basal cell carcinoma, dermatofibroma, melanoma, nevus, pigmented benign keratosis, seborrheic keratosis, squamous cell carcinoma, and vascular lesion.

3. Constraints and Data Management Analysis

The system design demonstrates efforts to mitigate data issues, but the evaluation phase presents a constraint:

- **Data Balancing Strategy:** During model development, a critical step was taken to address dataset imbalance. The training set was actively balanced to contain 1000 images for each of the nine classes, totaling 9000 images. This prevents the model from developing a bias toward over-predicting classes that were naturally more numerous in the original dataset.
- **Testing Set Limitation:** A constraint exists in the final evaluation environment. The test set, while totaling 118 images, is highly imbalanced for specific classes. Two classes—seborrheic keratosis and vascular lesion—are represented by only 3 images each.
- **Implication:** Due to the extremely low count of test samples for those two classes, the model's performance metrics (accuracy, precision, recall) for seborrheic keratosis and vascular lesion are less statistically robust and reliable compared to the performance metrics for the other seven classes, which have 16 images each in the test set.

2.2 Research Gap

The research gaps for the "Skin Cancer Detection using IIoT-based CNN" project can be categorized into four main areas: Model Performance and Complexity, Data Robustness and Generalization, IIoT Deployment and Efficiency, and Clinical Integration and Trust.

1. Model Performance and Complexity Gaps

1. **Insufficient Diagnostic Accuracy:** The reported validation accuracy (around 53.7%) is far below the threshold required for a clinically acceptable diagnostic tool, necessitating research into methods to significantly improve performance.
2. **Suboptimal Architecture Choice:** MobileNetV2, while efficient, may not possess the necessary feature extraction capability to differentiate subtle dermatological features. Exploration of more advanced architectures (e.g., Vision Transformers, EfficientNets) is needed.
3. **Lack of Comparative Analysis:** The project does not compare the MobileNetV2 fine-tuning results against other state-of-the-art models or against models trained from scratch, leaving the true efficiency of the chosen approach unverified.
4. **Limited Hyperparameter Optimization:** The project documentation suggests a standard training approach, but lacks detailed research into advanced learning rate schedules, novel optimizers, or specific regularization techniques that could boost model convergence and final performance.

II. Data Robustness and Generalization Gaps

5. **Critical Test Set Imbalance:** The severe underrepresentation of two lesion classes (seborrheic keratosis and vascular lesion, with only 3 test samples each) renders the model's performance on these specific classes statistically unreliable, indicating a gap in robust evaluation methodology.
6. **Lack of Rare-Case Handling Strategy:** Despite manual balancing of the training data, the model's true capability in handling the rare, yet critical, malignant lesions encountered in a natural, highly imbalanced clinical setting is unproven. Research into few-shot learning or synthetic data generation is needed.
7. **Single-Source Dataset Reliance:** The exclusive use of the HAM10000 dataset creates a generalization gap. The model is not validated against external, heterogeneous datasets captured with different cameras, lighting, and clinical protocols, limiting its real-world applicability.
8. **Absence of Patient-Level Generalization:** The system analysis does not confirm if the test data is patient-disjoint from the training data, which, if not addressed, can lead to an artificially inflated perception of the model's generalization capability.

III. IIoT Deployment and Efficiency Gaps

9. **Untested Edge Deployment Viability:** While the IIoT context is named, the project lacks a quantitative analysis of the model's performance when deployed on a specific, resource-constrained edge device (e.g., Raspberry Pi, specialized medical gateway).
10. **Missing Power Consumption Analysis:** There is no research detailing the energy efficiency of the model, a crucial metric for battery-powered IIoT medical devices where the model must execute tasks with minimal power drain.
11. **Lack of Latency Benchmarking:** Research is needed to benchmark the end-to-end inference latency on the target hardware to ensure the system meets the real-time requirements of clinical diagnostic assistance.
12. **Absence of Model Optimization Techniques:** The project does not include research into model optimization steps essential for edge computing, such as post-training quantization, network pruning, or knowledge distillation.

CHAPTER 3: METHODOLOGY

3.1 Data Collection

The data collection process for this skin cancer detection project is centered on acquiring and preparing a specific dataset for the multi-class image classification task. Since the project uses an established public dataset, the process focuses on procurement, standardization, and balancing rather than original clinical collection.

1. Data Source Procurement

The project utilizes the **HAM10000 dataset**, which is a collection of dermatoscopic images commonly used for machine learning research in skin lesion analysis. The data collection process begins by acquiring this dataset, which contains over 10,000 images of common pigmented skin lesions. The acquisition is done via digital download or access from the respective repository.

2. Initial Data Filtering and Standardization

Once procured, the dataset is initially processed to ensure all images conform to a standard for the deep learning model.

- **Image Resolution and Format:** All images are likely standardized to a uniform size (e.g., 224x224 pixels) to serve as consistent input for the MobileNetV2 architecture.
- **Classification Mapping:** The images are categorized based on their associated metadata into the nine target classes: Actinic keratosis, Basal cell carcinoma, Dermatofibroma, Melanoma, Nevus, Pigmented benign keratosis, Seborrheic keratosis, Squamous cell carcinoma, and Vascular lesion.

3. Training and Testing Set Partitioning

The total dataset is partitioned into three subsets to facilitate model training and evaluation:

- **Training Set:** Used to teach the model to recognize patterns and features associated with each lesion type.
- **Validation Set:** Used during the training phase to monitor the model's performance on unseen data and adjust hyperparameters to prevent overfitting.
- **Testing Set:** A completely separate and unseen subset used only once at the very end to provide a final, unbiased assessment of the model's performance.

4. Class Balancing Strategy

A critical step in the data collection and preparation process is addressing the inherent imbalance present in the original dataset.

- **Balancing Goal:** The primary objective was to ensure the training data did not bias the model toward the most common lesion types.
- **Implementation:** The researchers implemented an explicit balancing strategy, resulting in a perfectly balanced training set of **9000 images**—specifically **1000 images for each of the nine classes**. This likely involved oversampling the minority classes (duplicating or augmenting their images) and/or undersampling the majority classes.

5. Final Test Set Allocation (Evaluation Data)

After the training data is prepared, the remaining images are designated for evaluation.

- **Size:** The total size of the final test set is **118 images**.
- **Distribution:** While the training data was perfectly balanced, the test set was not, reflecting a natural distribution or a specific partitioning decision. The resulting distribution shows **16 images** for most classes, but only **3 images** each for the seborrheic keratosis and vascular lesion classes. This disparity means the data collection process resulted in limited, specific evaluation samples for these two categories.

3.2 Data Cleaning & Preprocessing

The preprocessing methodology focuses on standardization, normalization, and augmentation.

1. Image Standardization

All input images, regardless of their original dimensions, are resized to a uniform dimension. This ensures a consistent input shape for the CNN, which is a mandatory requirement for fixed-architecture models like MobileNetV2. This step removes variability in spatial dimensions.

2. Label Encoding and Categorization

The associated diagnostic labels (the nine skin lesion types) are converted into a numerical format suitable for machine learning. Since this is a multi-class classification problem, the labels are converted into a one-hot encoded format. This means each lesion type is represented as a binary vector (e.g., [0, 0, 0, 1, 0, 0, 0, 0, 0] for the fourth class, melanoma), which is necessary for training a model with a categorical cross-entropy loss function.

3. Pixel Value Normalization

The raw pixel intensity values of the images, typically ranging from 0 to 255, are scaled down. This normalization process converts the pixel values to a floating-point range, most commonly between 0 and 1. Normalization is essential as it stabilizes the training process, preventing large input values from disproportionately affecting the weights and biases of the neural network.

4. Data Augmentation

This is a critical preprocessing step used to artificially expand the size and diversity of the training set, which helps the model generalize better to unseen images and reduces overfitting. Common augmentation techniques applied to the image data likely include:

- **Rotation:** Randomly rotating images by small degrees.
- **Shifting:** Horizontally and vertically shifting the image content.
- **Flipping:** Mirroring the images along vertical or horizontal axes.
- **Zooming:** Randomly zooming in or out on the image.
- **Shearing:** Applying a shearing transformation to the image.

By generating these slightly modified copies, the model learns to recognize the lesions regardless of their specific position, orientation, or size variations captured during real-world imaging.

5. Batching and Buffering

The final step organizes the processed images into batches. The images are grouped together to enable efficient parallel processing on the GPU hardware (a T4 GPU was used in the training environment). The data pipeline is set up to load these batches asynchronously, ensuring that the GPU is constantly fed data without delays, optimizing the overall training speed.

3.3 Data Modeling

1. **Architecture:** The model uses **MobileNetV2**, an efficient and lightweight Convolutional Neural Network (CNN), chosen specifically for its suitability for resource-constrained IIoT edge deployment.
2. **Technique:** The project implements **transfer learning**. The pre-trained MobileNetV2 base model is loaded, and its weights are frozen to preserve general image features.
3. **Customization:** A new classification head with **nine output nodes** is added to the frozen base to match the nine skin lesion classes.
4. **Configuration:** The model is compiled using the **Categorical Cross-Entropy** loss function, which is standard for multi-class problems with one-hot encoded labels.

CHAPTER 4: SYSTEM DESIGN

4.1 Tools & Technologies Used

The tools and technologies used in this Skin Cancer Detection project span the hardware, programming, and software framework layers, reflecting a typical modern deep learning implementation in an Industrial Internet of Things (IIoT) context.

1. Hardware and Computing Environment

- **Graphics Processing Unit (GPU):** The project utilized a **T4 GPU** for training, as indicated in the notebook metadata. This is a powerful, cloud-based processor essential for accelerating the computationally intensive operations of training deep Convolutional Neural Networks (CNNs).
- **Execution Environment:** The code was executed within a platform like **Google Colab** (suggested by the notebook format and GPU type), providing a cloud-hosted environment necessary for the development and training process.

2. Core Programming Languages and Libraries

- **Programming Language: Python** serves as the primary language for all scripting, data handling, and model construction.
- **Deep Learning Framework: TensorFlow** (abbreviated as tf) is the foundational deep learning library used. This framework provides the core functionalities for defining, training, and saving the neural network model.
- **High-Level API: Keras**, which runs on top of TensorFlow, is used to define and structure the neural network. Keras simplifies the process of model building, especially for complex architectures like MobileNetV2, by providing an intuitive, high-level API.

3. Model Architecture and Components

- **Base Model: MobileNetV2** is the specific pre-trained CNN architecture employed. This model is a specialized network designed for efficiency and high performance on mobile and edge devices, aligning perfectly with the IIoT context of the project.
- **Utility Libraries:** Standard Python libraries are utilized for essential data and file management tasks:
 - **NumPy:** Used for high-performance numerical operations, particularly for handling arrays and matrices, which are the fundamental data structures for images and model weights.
 - **Pandas:** Used for data manipulation and analysis, likely employed to load and process the metadata associated with the HAM10000 dataset.
 - **Matplotlib / Seaborn:** Used for data visualization, specifically for plotting training and validation metrics (accuracy, loss) and potentially visualizing the images and the confusion matrix.

CHAPTER 5: RESULTS & ANALYSIS

5.1 Insights / Observations

The detailed insights and observations derived from the project data and execution logs focus on the data preparation, model training progression, and the implications of the test set distribution.

1. Data Preparation and Integrity Insights

- **Training Data Balancing Success:** The project successfully implemented a crucial data cleaning step by ensuring a perfectly balanced training dataset. There are exactly **1000 images** for each of the nine skin lesion classes, resulting in a total training set of **9000 images**. This active balancing prevents the model from developing a bias toward naturally more common lesion types in the initial data.
- **Test Data Imbalance Failure:** A critical observation is the severe imbalance in the evaluation set. The total test image count is only **118**. While most of the nine classes have a count of **16 images** each, two specific classes—**seborrheic keratosis** and **vascular lesion**—are drastically underrepresented, each having only **3 images**.
- **Evaluation Reliability Issue:** The low count of 3 images for the two under-represented classes in the test set means that the model's performance metrics (like accuracy) calculated for these two specific lesions will be statistically unreliable and subject to high variance. This is a major limitation of the current evaluation design.

2. Model Training and Performance Insights

- **Learning Progression:** The training process, monitored over 25 epochs, shows that the model is actively learning.
- **Initial Performance:** The model started with a low validation accuracy of approximately **31.32 percent** and a validation loss of **1.7752** in Epoch 1.
- **Accuracy Improvement:** The validation accuracy shows a general upward trend, stabilizing and reaching **53.69 percent** by Epoch 17. The validation loss simultaneously decreased to **1.3921** by Epoch 17.
- **Performance Ceiling:** The observed maximum validation accuracy, around 53.7 percent, is low for a multi-class medical diagnosis system. This suggests the model, in its current state, is not yet clinically viable and requires significant architecture changes or hyperparameter tuning.

3. Execution Environment Insights

- **Hardware Confirmation:** The project was executed using a **T4 GPU**. This indicates that the training phase leveraged high-performance computing resources necessary for training deep neural networks.
- **IIoT Implication:** The choice of the MobileNetV2 architecture, combined with the use of a powerful T4 GPU for training, highlights the practical design goal: developing a model on high-end hardware that is efficient enough for later deployment on less powerful, dedicated IIoT devices.

CHAPTER 6: CONCLUSION & FUTURE SCOPE

6.1 Conclusion

The project successfully implemented an Industrial Internet of Things (IIoT)-based system for **nine-class skin cancer detection** using a deep learning methodology, but its current state is best characterized as a promising proof-of-concept that requires significant further development before achieving clinical utility.

The **objective** of creating an automated diagnostic aid for nine distinct skin lesions was technically met by constructing a classifier based on **transfer learning** with the **MobileNetV2** architecture. The project demonstrated sound data preparation practices by successfully balancing the **9000-image training dataset** (1000 images per class), which is crucial for mitigating class bias during the learning phase. The training process itself confirmed that the model was actively learning, with validation accuracy rising from approximately 31% to around 53.7 % over the observed epochs.

However, the final conclusion highlights several critical limitations:

Limitations and Path Forward

1. **Insufficient Clinical Performance:** The final observed validation accuracy of around 53.7 % is far too low for a medical diagnostic tool where high confidence and reliability are paramount. The model, in its current form, is not accurate enough to be deployed or relied upon by clinicians.
2. **Evaluation Reliability Issues:** The evaluation methodology is constrained by a highly imbalanced test set, particularly the presence of only 3 images each for seborrheic keratosis and vascular lesion. This makes the system's performance metrics on these specific classes unreliable, requiring a larger, more balanced, and externally validated test set.
3. **IIoT Deployment Gap:** While the MobileNetV2 architecture is theoretically suited for IIoT edge devices due to its efficiency, the project lacks quantitative research (e.g., latency, power consumption, model quantization) to prove its practical, real-time deployability on specific low-power hardware.
4. **Lack of Transparency:** The absence of Explainable AI (XAI) techniques presents a major hurdle for clinical adoption. Without tools to show *why* the model made a certain diagnosis, the system remains a "black box" that medical professionals are unlikely to trust or integrate into patient care protocols.

In summary, the project established a strong foundation by defining the problem, selecting an efficient architecture, and performing careful training data balancing. The next phase of development must pivot from a focus on foundational methodology to rigorous performance enhancement, robust external validation, optimization for true edge computing efficiency, and the integration of clinically necessary interpretability features to close the gap between the prototype and a viable IIoT-enabled diagnostic product.

6.2 Future Scope

1. Advanced Model Research and Performance Enhancement

- **Architecture Migration:** Future work must move beyond MobileNetV2. Research should focus on implementing and benchmarking more powerful and contemporary CNN architectures or Vision Transformers (e.g., EfficientNet, Swin Transformer) to significantly raise diagnostic accuracy above the current $\approx 53.7\%$ mark.
- **Addressing Data Rarity:** Implement advanced learning techniques such as **Few-Shot Learning** or **Meta-Learning** to ensure the model maintains high predictive accuracy for rare and difficult-to-diagnose lesion classes, even those with limited training samples.
- **Multi-Modal Fusion:** Explore fusing image data with clinical metadata (age, gender, lesion site, patient history) if available. Integrating structured data with visual features can potentially boost the model's diagnostic power beyond image-only classification.
- **Uncertainty Quantification:** Research and implement methods (e.g., Bayesian Neural Networks or ensemble techniques) to provide a reliable measure of prediction confidence. This is crucial for guiding clinicians on when to trust the AI diagnosis and when to perform manual verification.

2. IIoT and Edge Deployment Optimization

- **Model Optimization and Compression:** Future research needs to focus heavily on model compression techniques, including **pruning** (removing redundant weights) and **quantization** (reducing the precision of weights, e.g., from 32-bit to 8-bit integers). This is vital to shrink the model size and speed up inference time on low-power IIoT hardware.
- **Hardware Benchmarking:** Conduct explicit testing of the optimized model on specific IIoT platforms (e.g., custom medical edge processors, NVIDIA Jetson devices). This research should measure and report real-world metrics: inference latency, power consumption (in watts), and frames per second (FPS) to guarantee real-time clinical operation.
- **Secure Data Pipeline Integration:** Develop and implement secure, low-latency communication protocols (e.g., MQTT or CoAP) for transmitting images from the dermatoscope to the edge device and sending the classification results back to a central Electronic Health Record (EHR) system, ensuring HIPAA/GDPR compliance.

3. Clinical Validation and Integration

- **External Validation:** The model must be validated on large, diverse, and independent datasets collected from different clinics and using various imaging equipment to ensure generalizability and prevent overfitting to the specific features of the HAM10000 dataset.
- **Explainable AI (XAI) Implementation:** Integrate XAI techniques, such as **Grad-CAM** (Gradient-weighted Class Activation Mapping), to generate heatmaps that highlight the specific visual regions the model uses to make its diagnosis. This provides transparency and builds trust with clinicians.
- **Clinical Trial Simulation:** Conduct simulation studies or pilot clinical trials with dermatologists to evaluate the system's utility, workflow impact, and diagnostic improvement compared to human-only diagnosis.

4. Commercialization and Regulatory Scope

- **Certification Research:** Initiate research and groundwork necessary for obtaining regulatory approval (e.g., FDA clearance in the US or CE marking in Europe). This involves stringent documentation, quality management systems, and proof of clinical safety and effectiveness.
- **Product Feature Expansion:** Scope the addition of features such as malignancy tracking, longitudinal monitoring of lesions, and integration of a teledermatology platform where difficult cases can be flagged for human review.
- **Cost-Benefit Analysis:** Perform a detailed analysis of the system's economic impact, quantifying the reduction in misdiagnoses, time saved per patient, and overall healthcare cost savings to justify its adoption in clinical settings.